

TM

VOXX



# *Dive into* **Spark** *Streaming* *by*

Gerard Maas

Data Processing Team Lead




[gerard.maas@virdata.com](mailto:gerard.maas@virdata.com)



@maasg

 virdata



Tweet few keywords about your interests and experience.

Use hashtag “#streamvoxx”

Compose new Tweet

#streamvoxx #spark #scala #bigdata #Devoxx

Media

Location

Poll

97

Tweet









100TB

5MB



100TB

5MB / second

# Agenda

Lightning fast Spark recap

What is Spark Streaming?

Programming Model

Demo 1

Execution Model

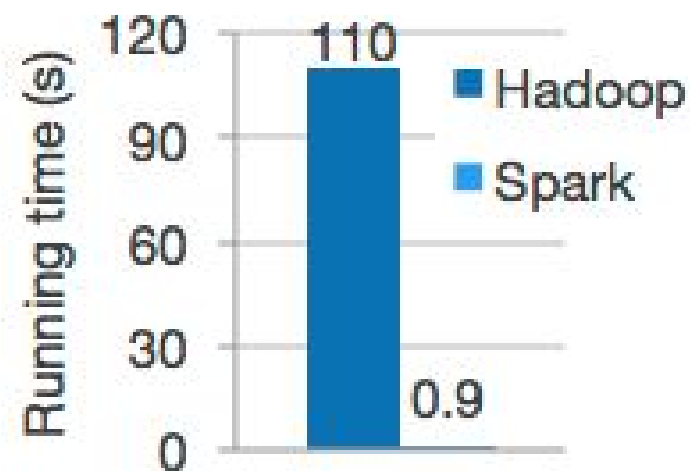
Demo 2

Resources

Q/A

# Apache Spark

Spark is a fast and general engine for large-scale distributed data processing.



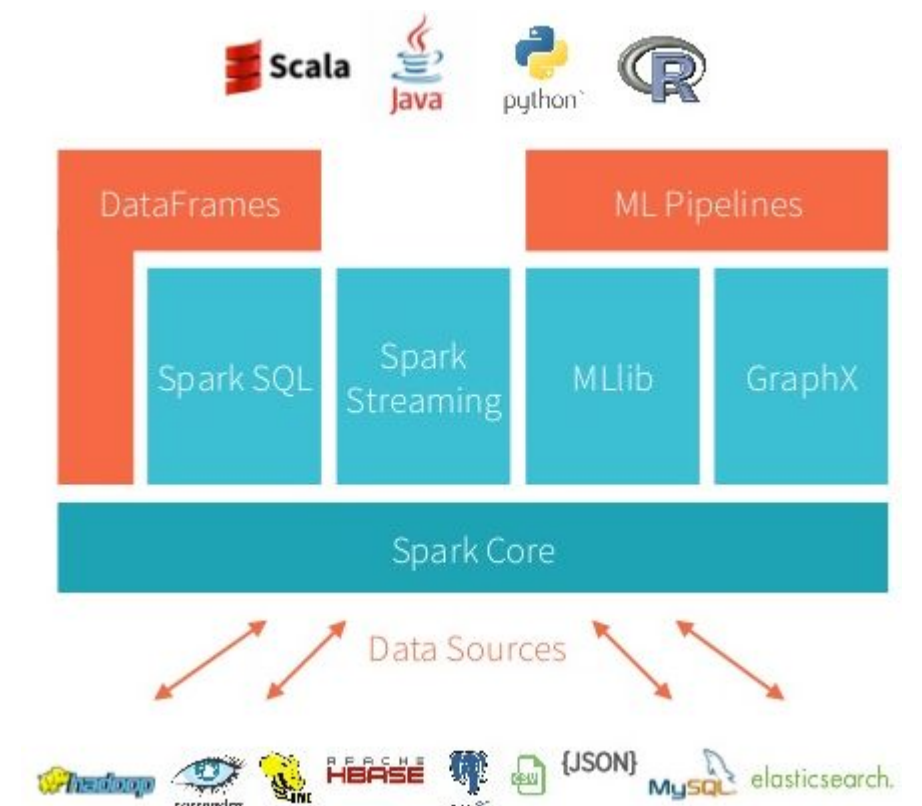
**Fast**

```
val file = spark.textFile("hdfs://...")

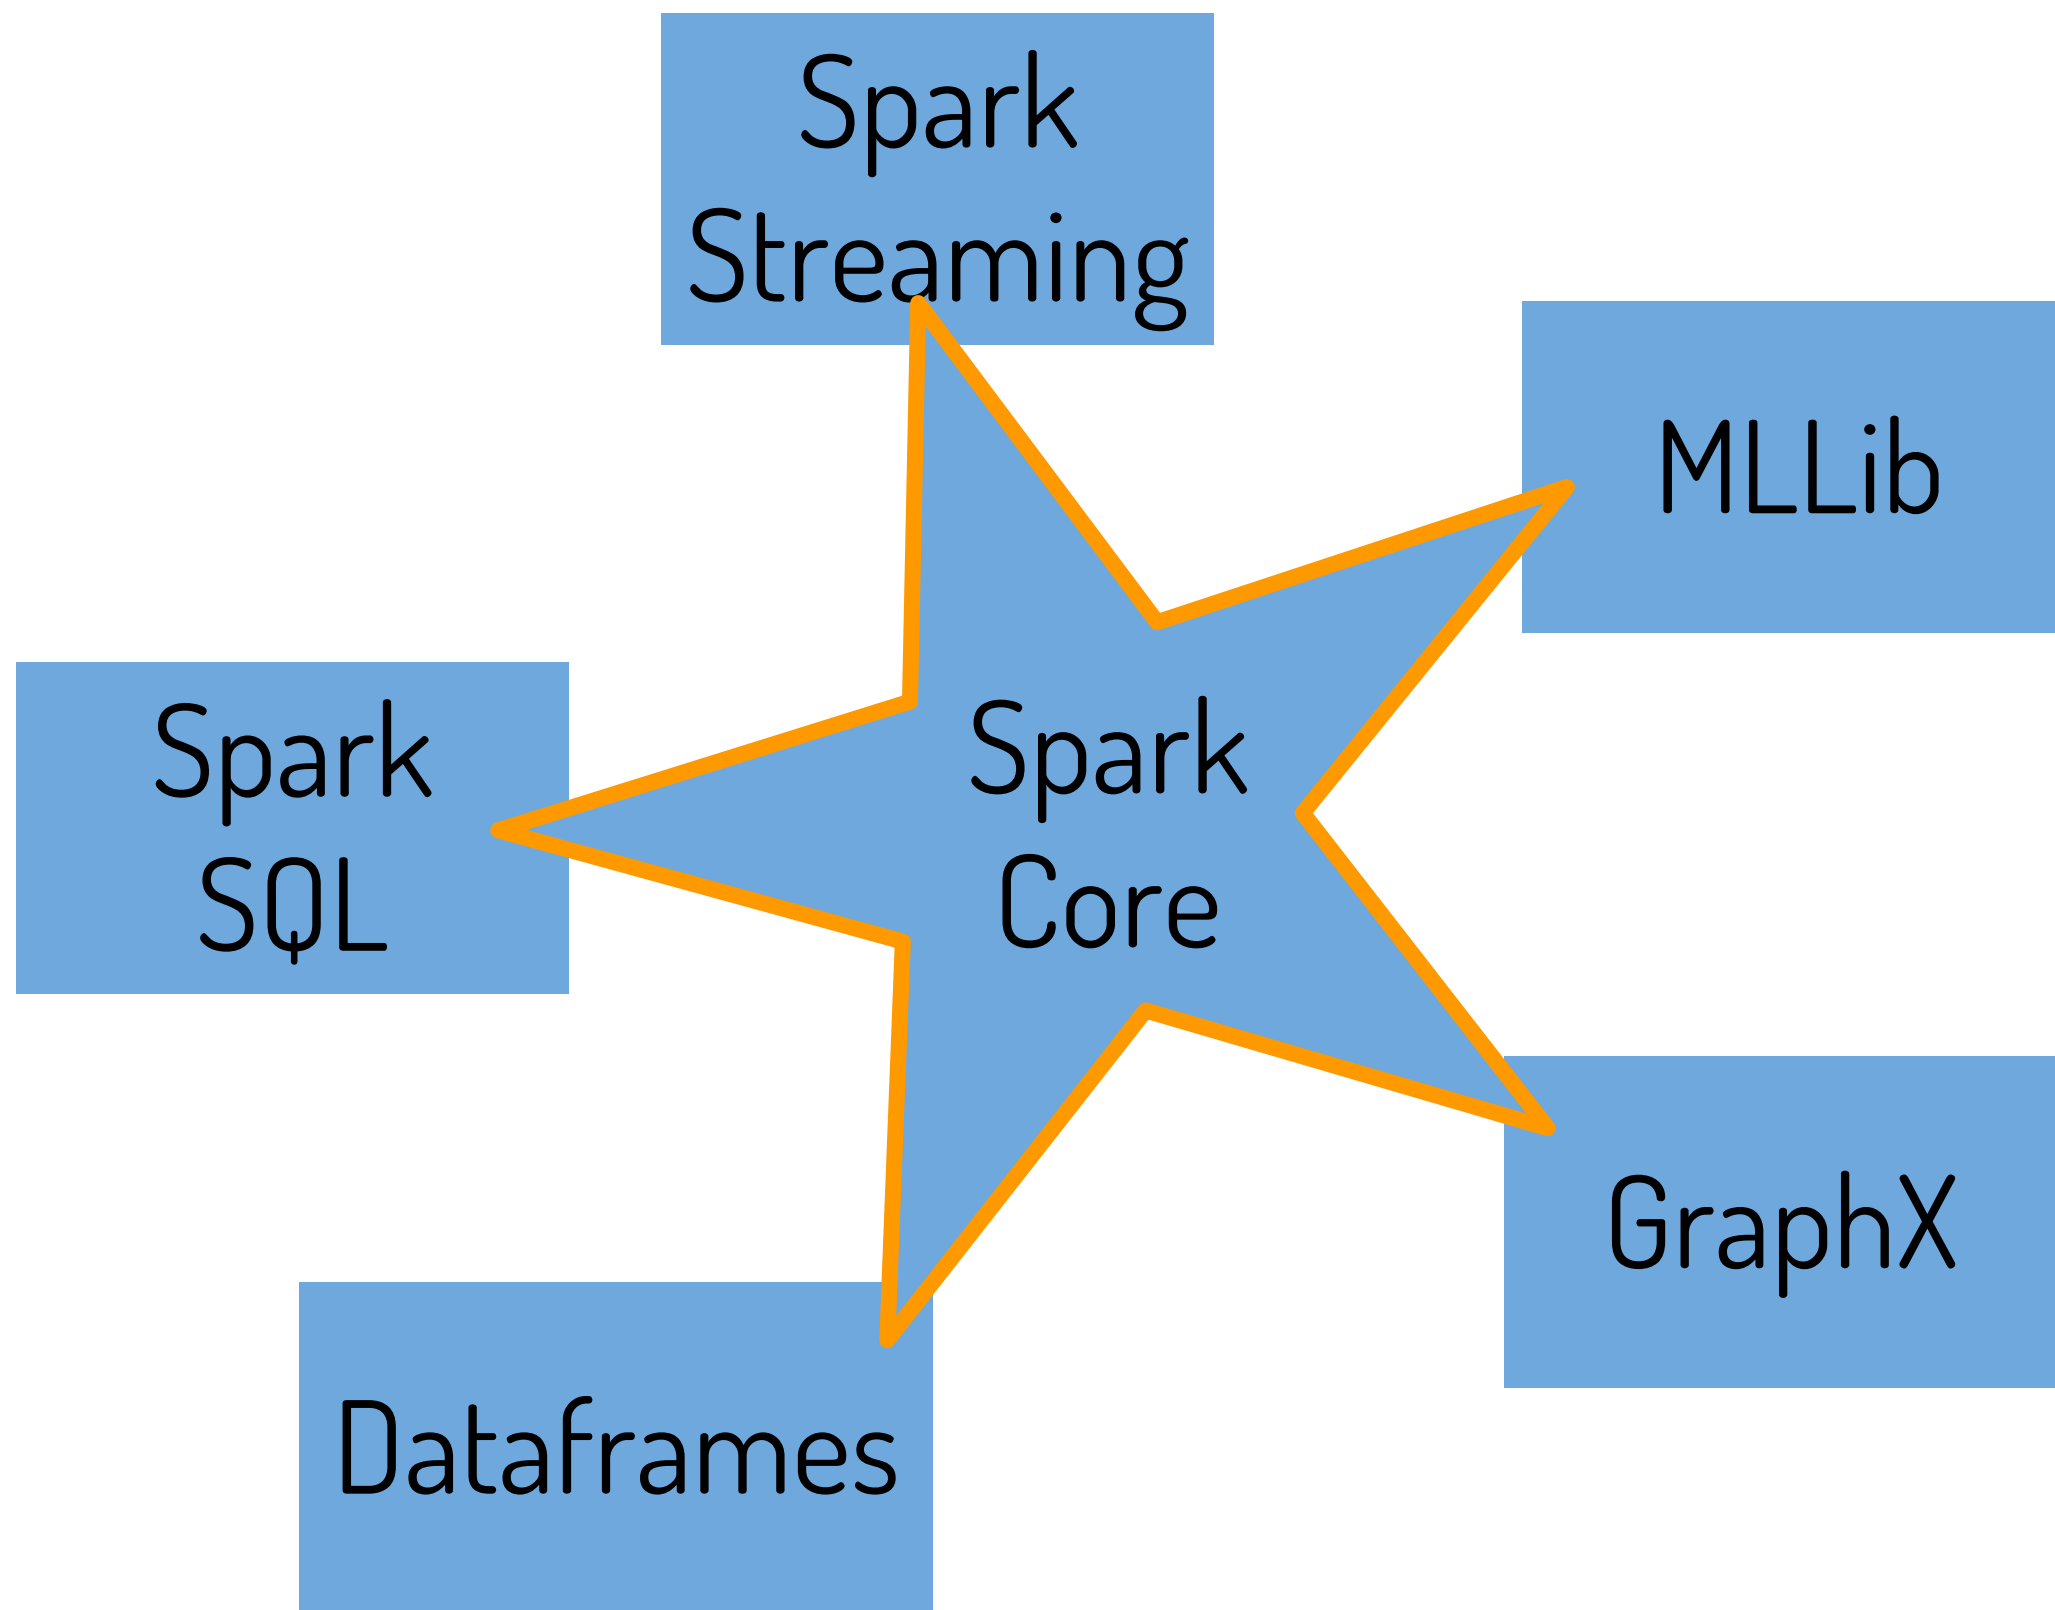
val counts = file.flatMap(line => line.
split(" "))
.map(word => (word, 1))
.reduceByKey(_ + _)

counts.saveAsTextFile("hdfs://...")
```

**Functional**



**Growing  
Ecosystem**







# RDDs

**Express computations in terms of transformations and actions on a distributed data set.**

Spark Core Concept: RDD => **R**esilient **D**istributed **D**ataset

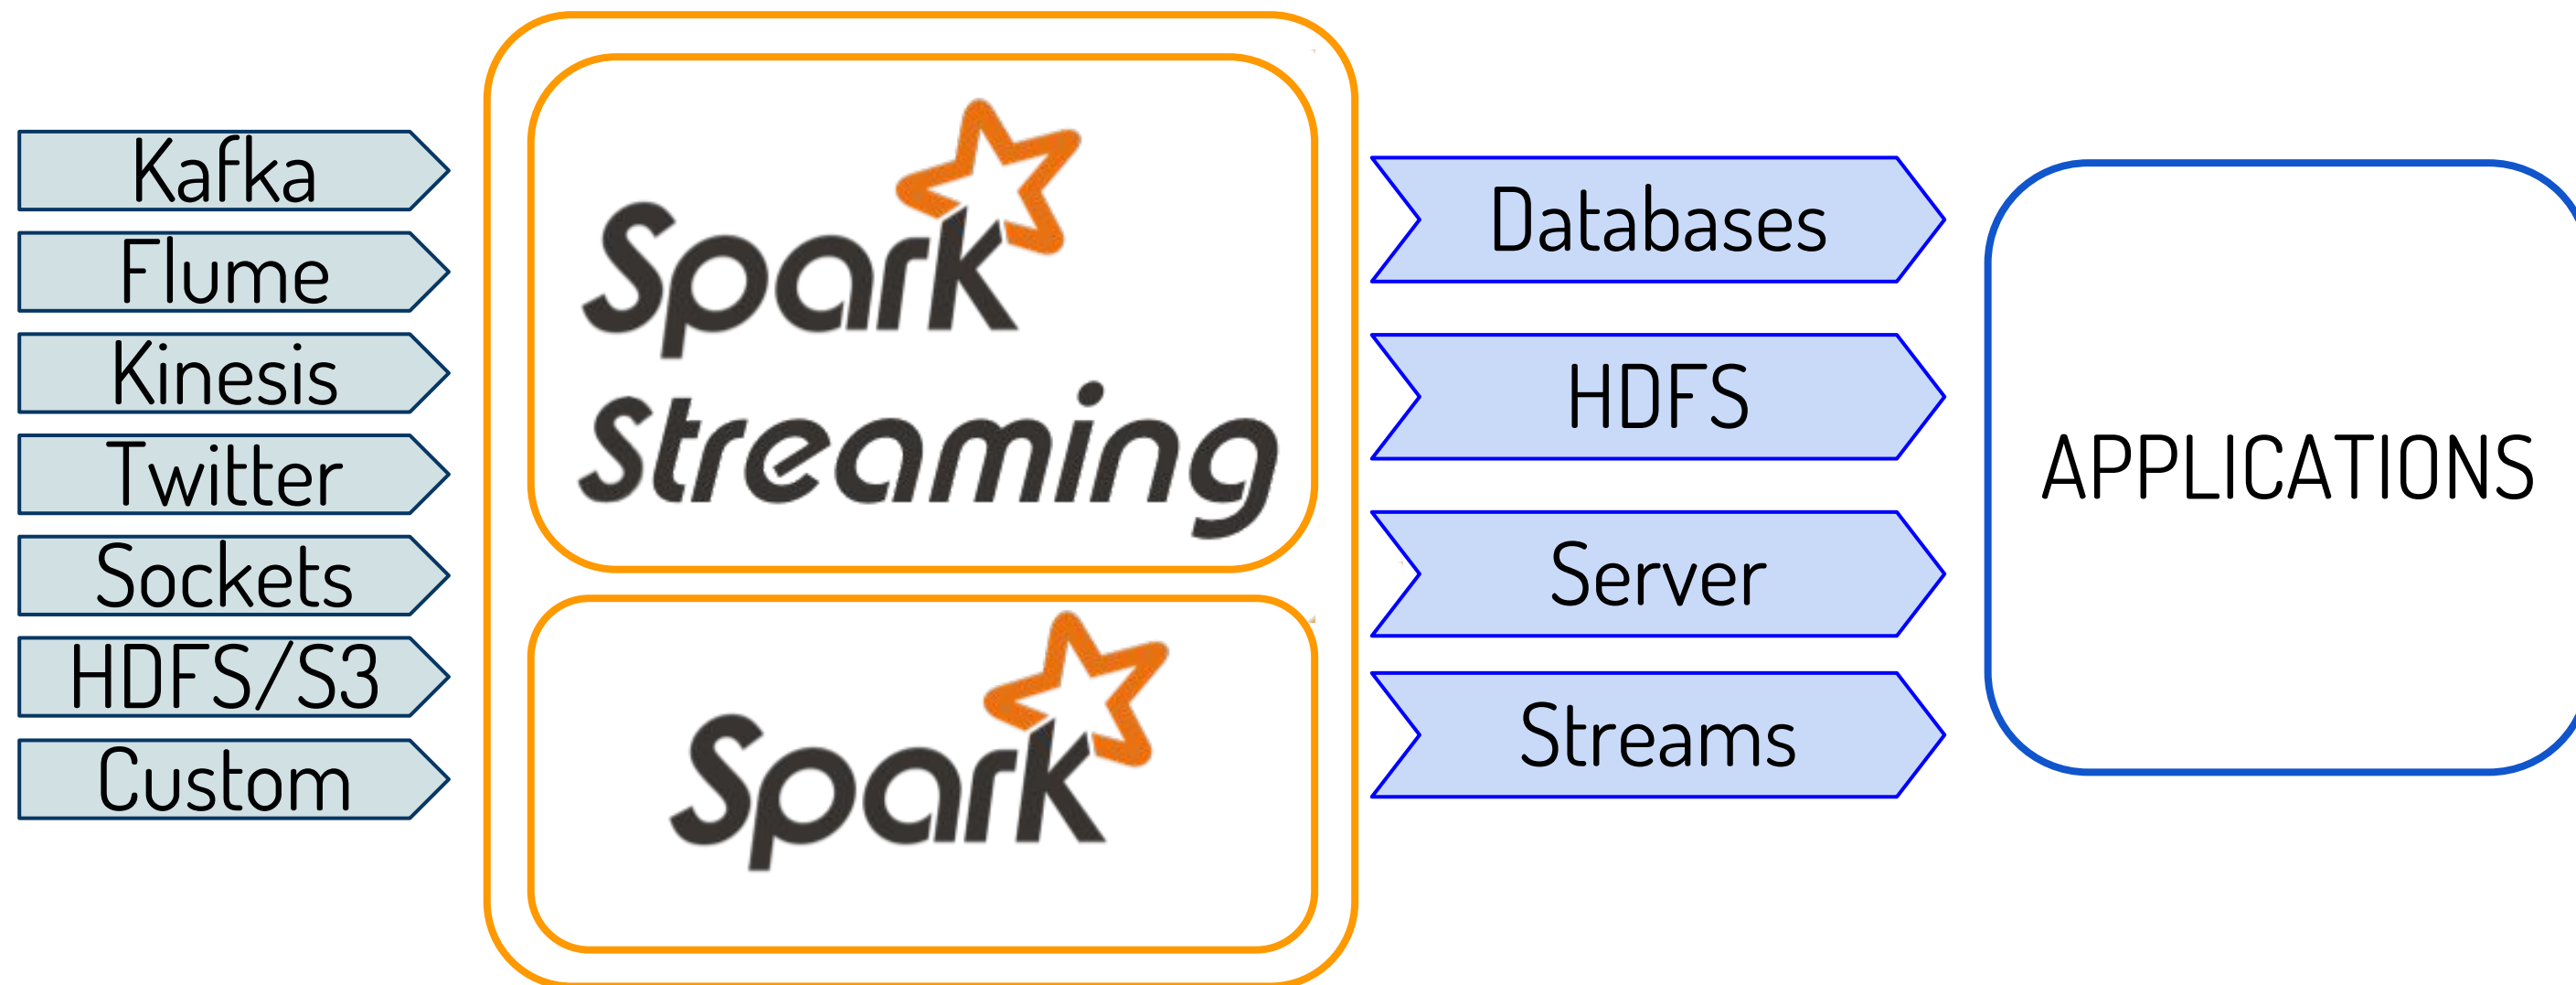
Think of an **RDD** as an **immutable**, **distributed** collection of **objects**

- **R**esilient => Can be reconstructed in case of failure
- **D**istributed => Transformations are parallelizable operations
- **D**ataset => Data loaded and partitioned across cluster nodes (executors)


RDDs are memory-intensive. **Caching** behavior is controllable.

# Spark Streaming

Scalable, fault-tolerant stream processing system




# Spark Streaming

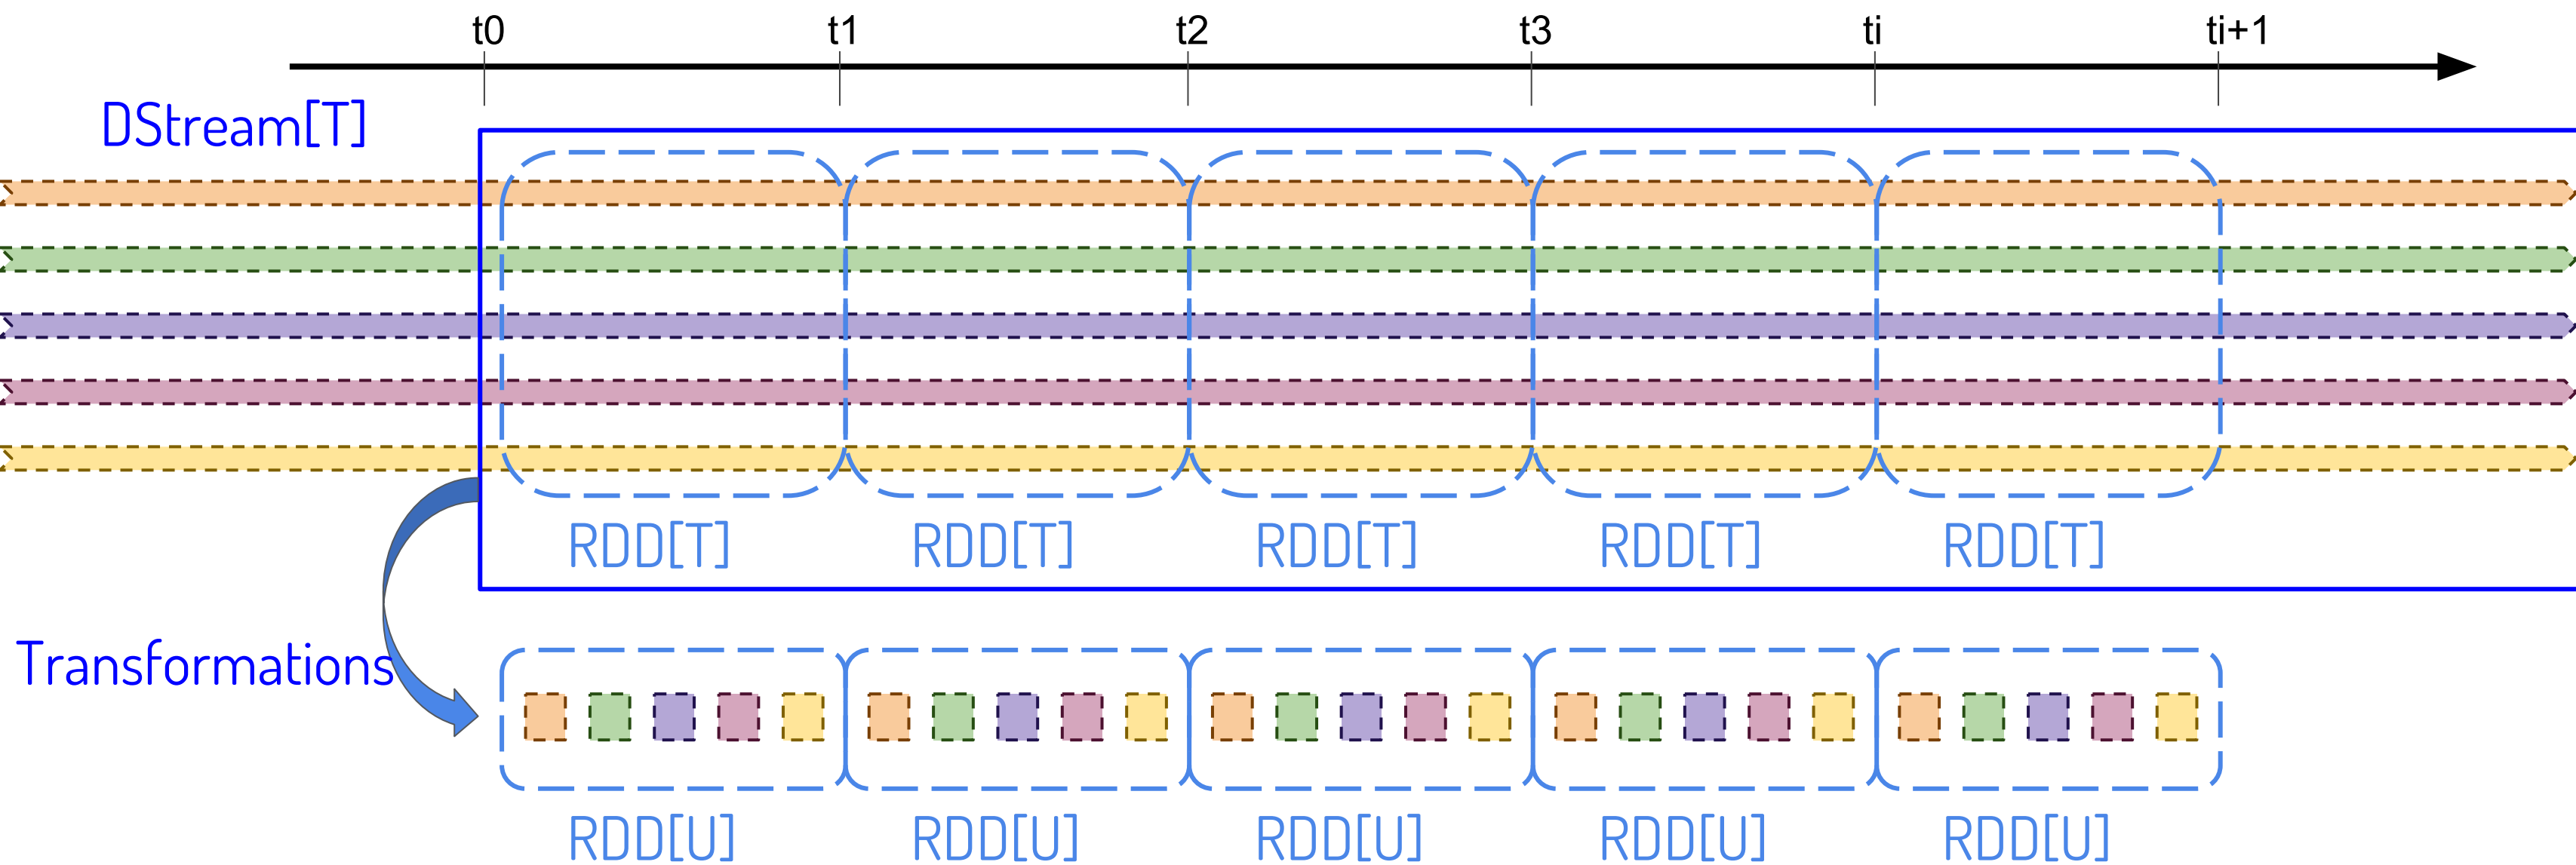




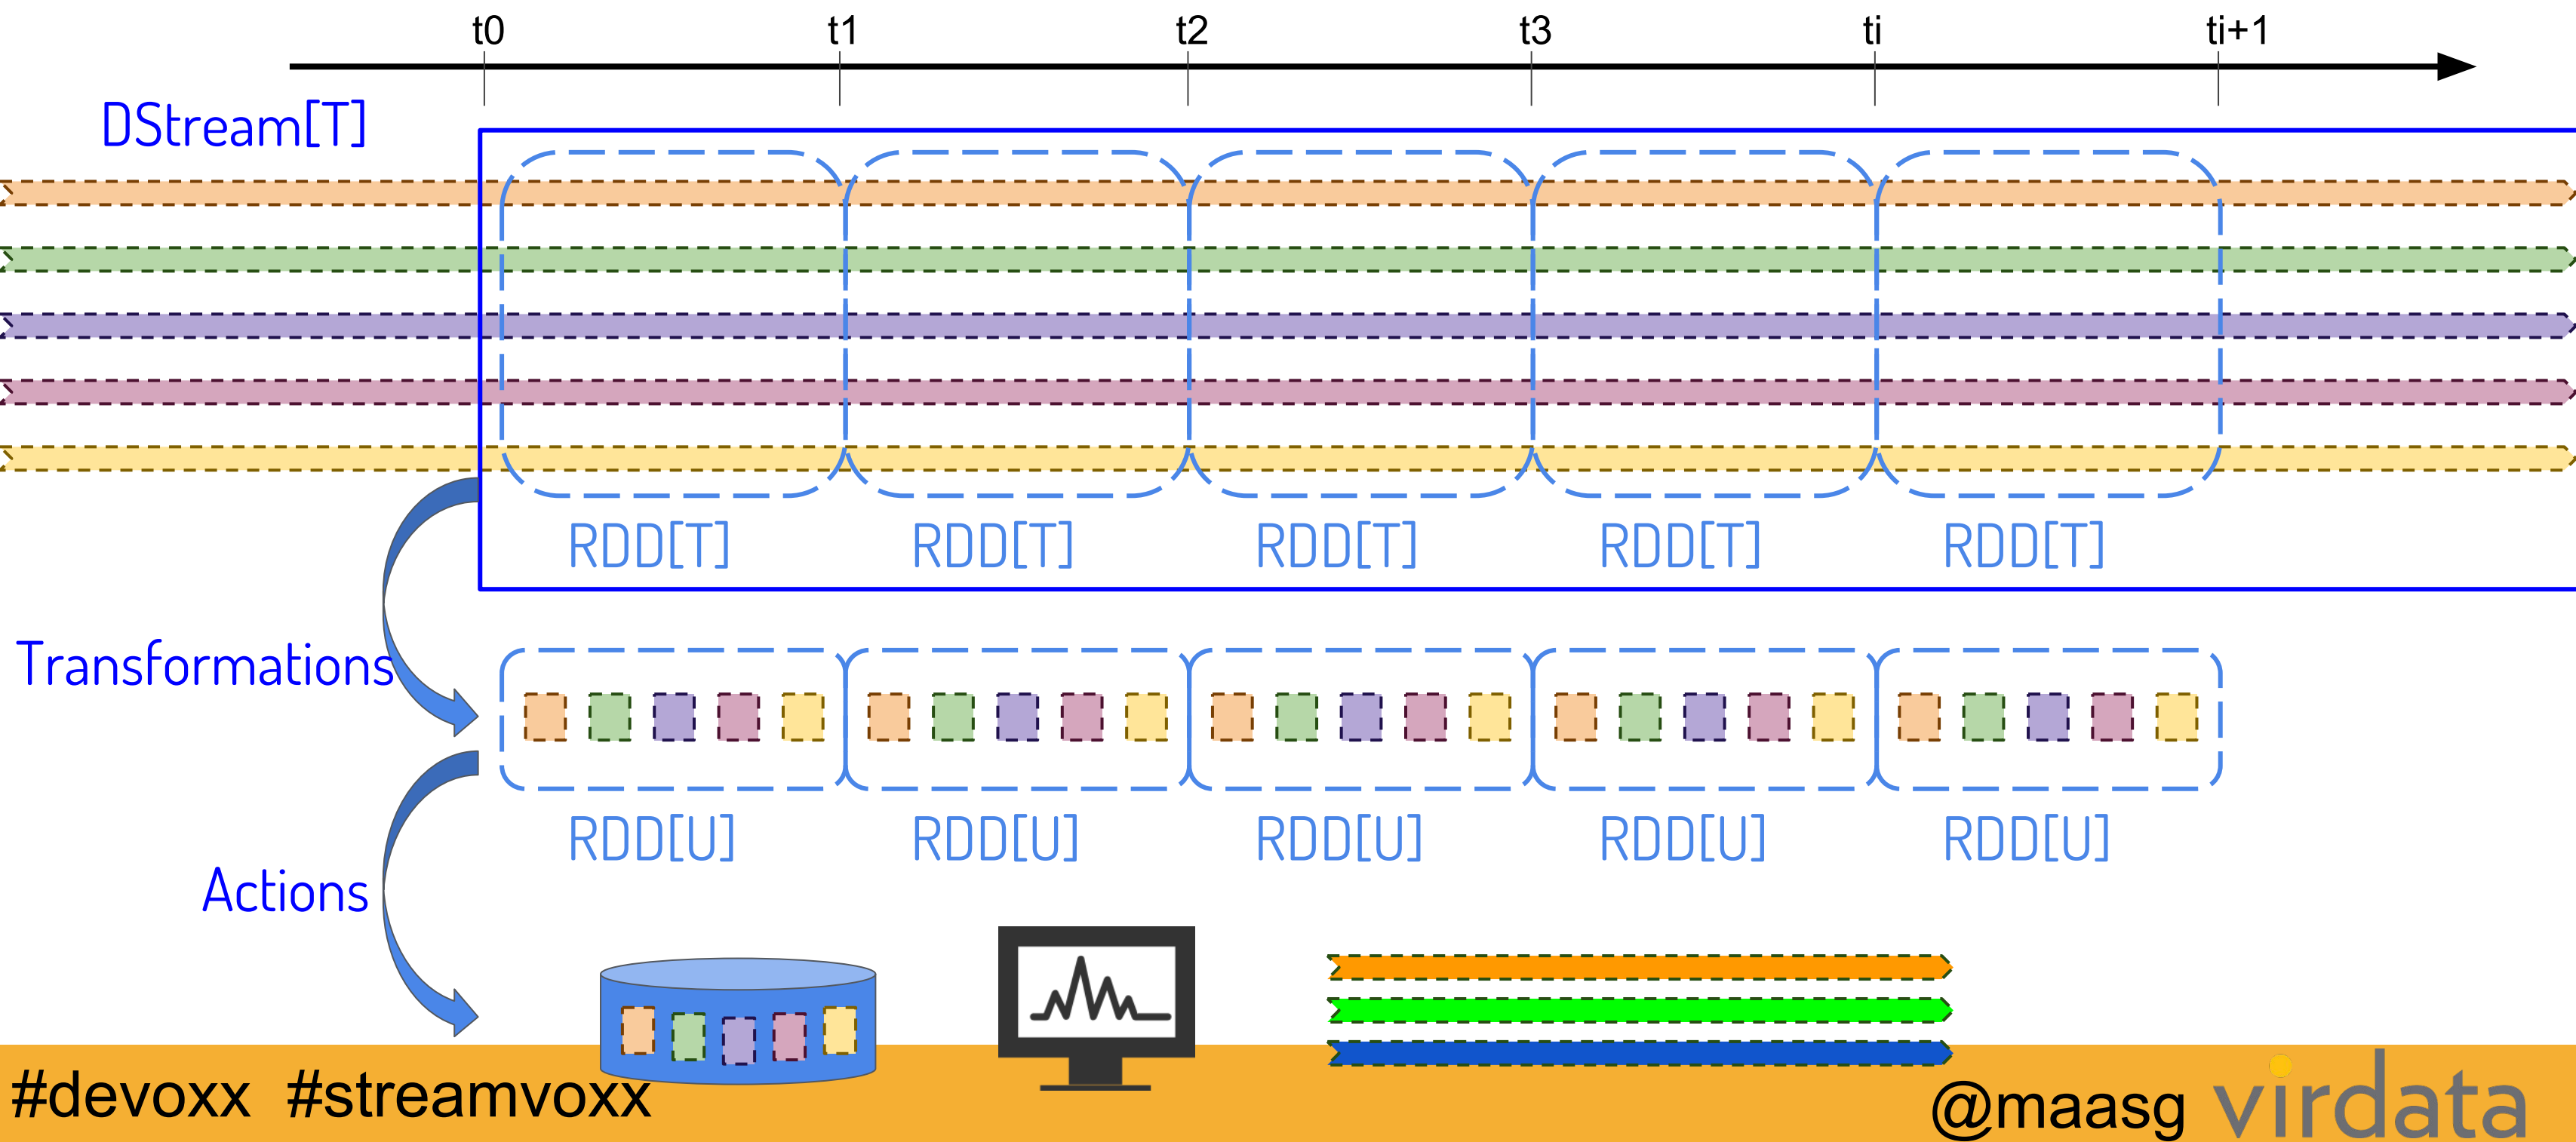
# Spark Streaming



# Spark Streaming




# Spark Streaming






# Transformations


map,  
flatmap,  
filter



count,  
reduce,  
countByValue,  
reduceByKey




union,  
join  
cogroup



# Transformations


transform



```
val iotDstream = MQTTUtils.createStream(...)
val devicePriority = sparkContext.cassandraTable(...)
val prioritizedDStream = iotDstream.transform{rdd =>
  rdd.map(d => (d.id, d)).join(devicePriority)
}
```

# Transformations


updateStateByKey





# Actions

print




Time: 1459875469000 ms


data1

data2

saveAsTextFiles,  
saveAsObjectFiles,  
saveAsHadoopFiles



foreachRDD



Actions – foreachRDD

```
dstream.foreachRDD{rdd =>
```

Spark SQL

MLLib

Dataframes

GraphX

Databases

...

```
}
```

# Actions – foreachRDD Usage

```
dstream.foreachRDD{rdd =>
    rdd.cache()
    val alternatives = restServer.get("/v1/alternatives").toSet
    alternatives.foreach{alternative =>
        val byAlternative = rdd.filter(element => element.kind == alternative)
        val asRecords = byAlternative.map(element => asRecord(element))
        val conn = DB.connect(server)
        asRecords.foreachPartition{partition =>
            partition.foreach(element => conn.insert(element))
        }
    }
    rdd.unpersist(true)
}
```



# Actions – foreachRDD Usage

```
dstream.foreachRDD{rdd =>
  rdd.cache()
  val alternatives = restServer.get("/v1/alternatives").toSet
  alternatives.foreach{alternative =>
    val byAlternative = rdd.filter(element => element.kind == alternative)
    val asRecords = byAlternative.map(element => asRecord(element))
    val conn = DB.connect(server)
    asRecords.foreachPartition{partition =>
      partition.foreach(element => conn.insert(element))
    }
  }
  rdd.unpersist(true)
}
```

Executes on the Driver

Executes on the Workers

## Actions – foreachRDD Usage


```
dstream.foreachRDD{rdd =>
  rdd.cache()
  val alternatives = restServer.get("/v1/alternatives").toSet
  alternatives.foreach{alternative =>
    val byAlternative = rdd.filter(element => element.kind == alternative)
    val asRecords = byAlternative.map(element => asRecord(element))
    asRecords.foreachPartition{partition =>
      val conn = DB.connect(server)
      partition.foreach(element => conn.insert(element))
    }
  }
  rdd.unpersist(true)
}
```

Executes on the Driver

Executes on the Workers

# Windows - Sliding


```
dstream.window(windowLength = 6, slideInterval = 3)
```





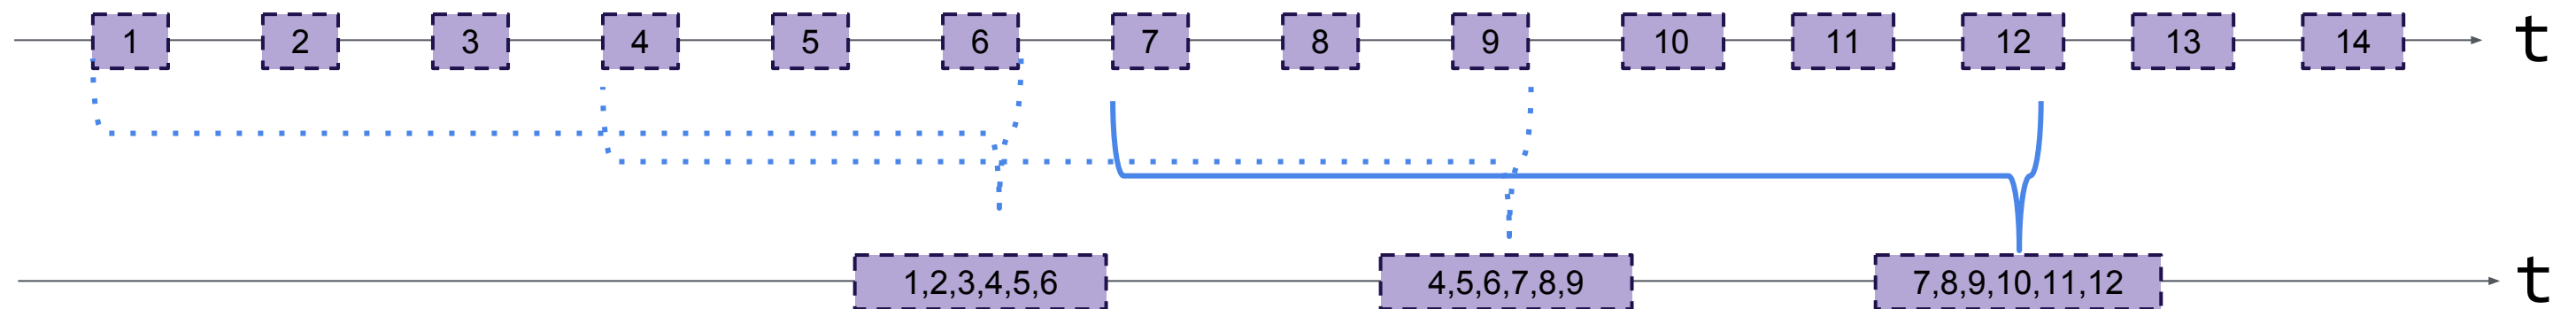
# Windows - Sliding

```
dstream.window(windowLength = 6, slideInterval = 3)
```



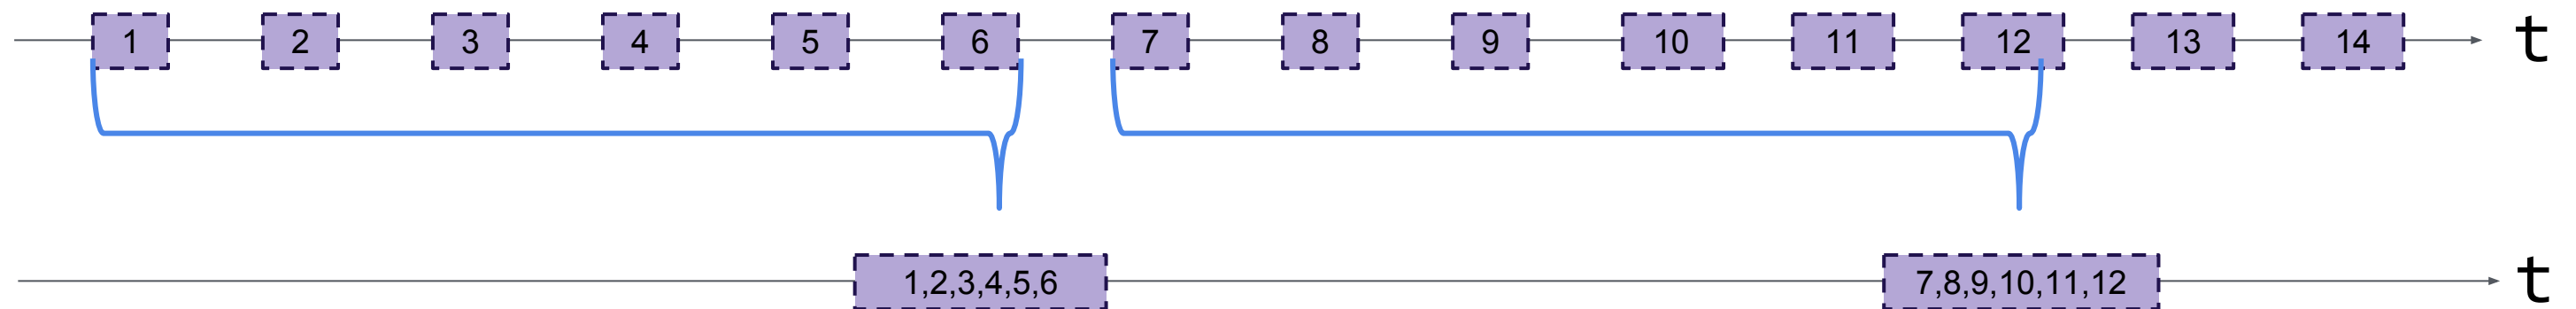
# Windows - Sliding

```
dstream.window(windowLength = 6, slideInterval = 3)
```



# Windows - Non-Overlapping

```
dstream.window(windowLength = 6, slideInterval = 6)
```




# Windows - Operations

window,  
countByWindow,  
reduceByWindow,  
reduceByKeyAndWindow,  
countByValueAndWindow

# Windows - Inverse Function Optimization

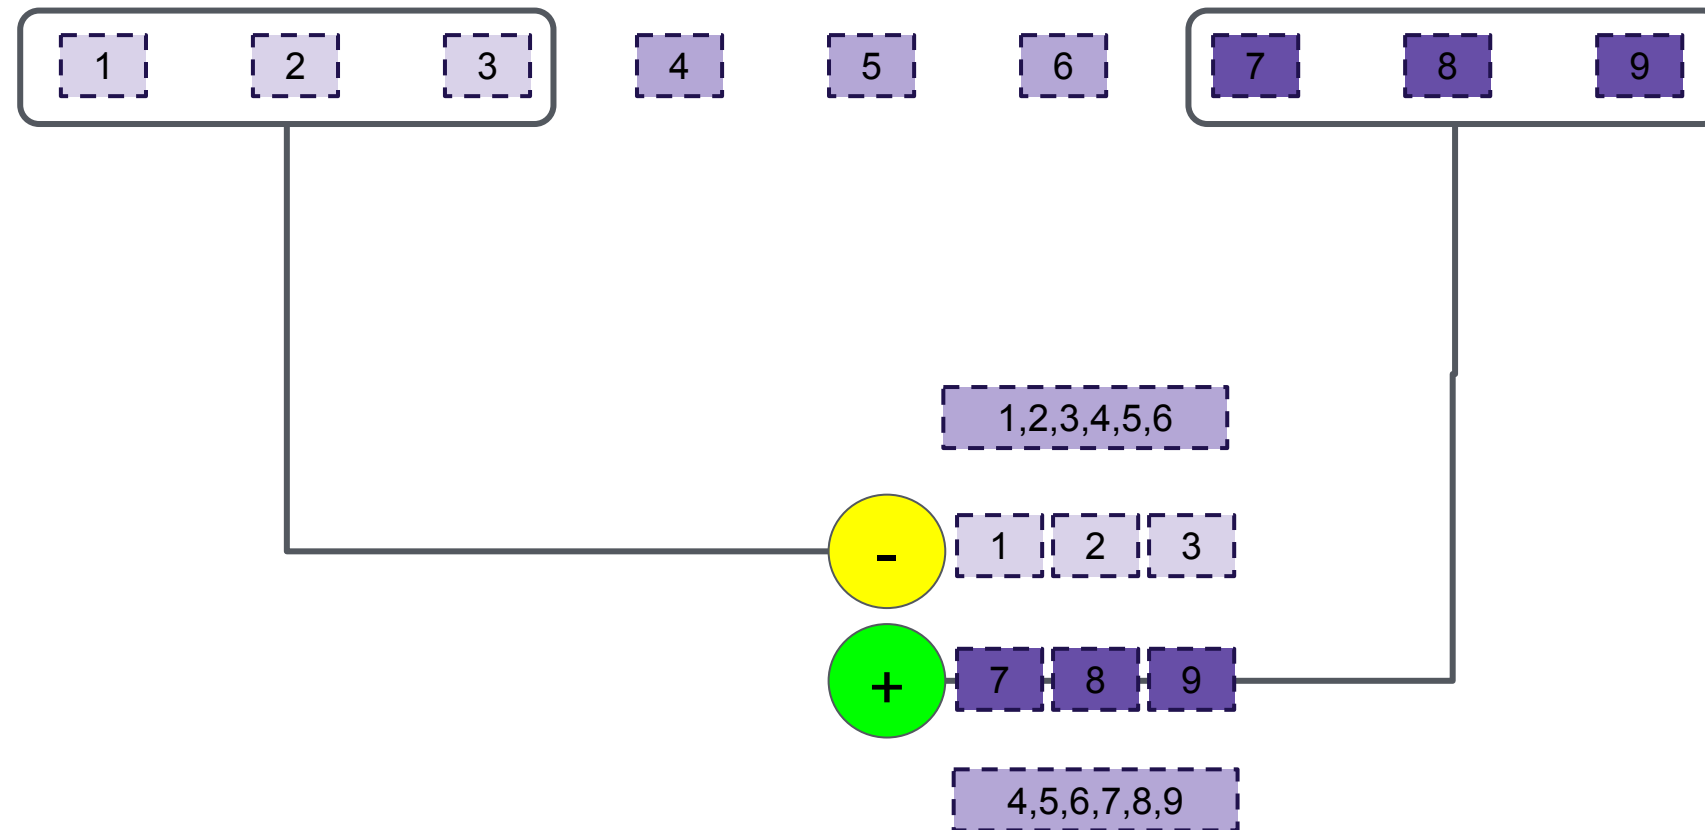
   
`reduceByKeyAndWindow(func, invFunc, windowLength, slideInterval, [numTasks])`





# Windows- Inverse Function Optimization

   
`reduceByKeyAndWindow(func, invFunc, windowLength, slideInterval, [numTasks])`





# Demo 1

Anatomy of an  
Spark Streaming  
Application



Ready to dive in?










# Deployment Options

Local




`spark.master=local[*]`

Standalone Cluster



`spark.master=spark://host:port`

Using a Cluster Manager



`spark.master=mesos://host:port`






# Deployment Options

Local




`spark.master=local[*]`

Standalone Cluster



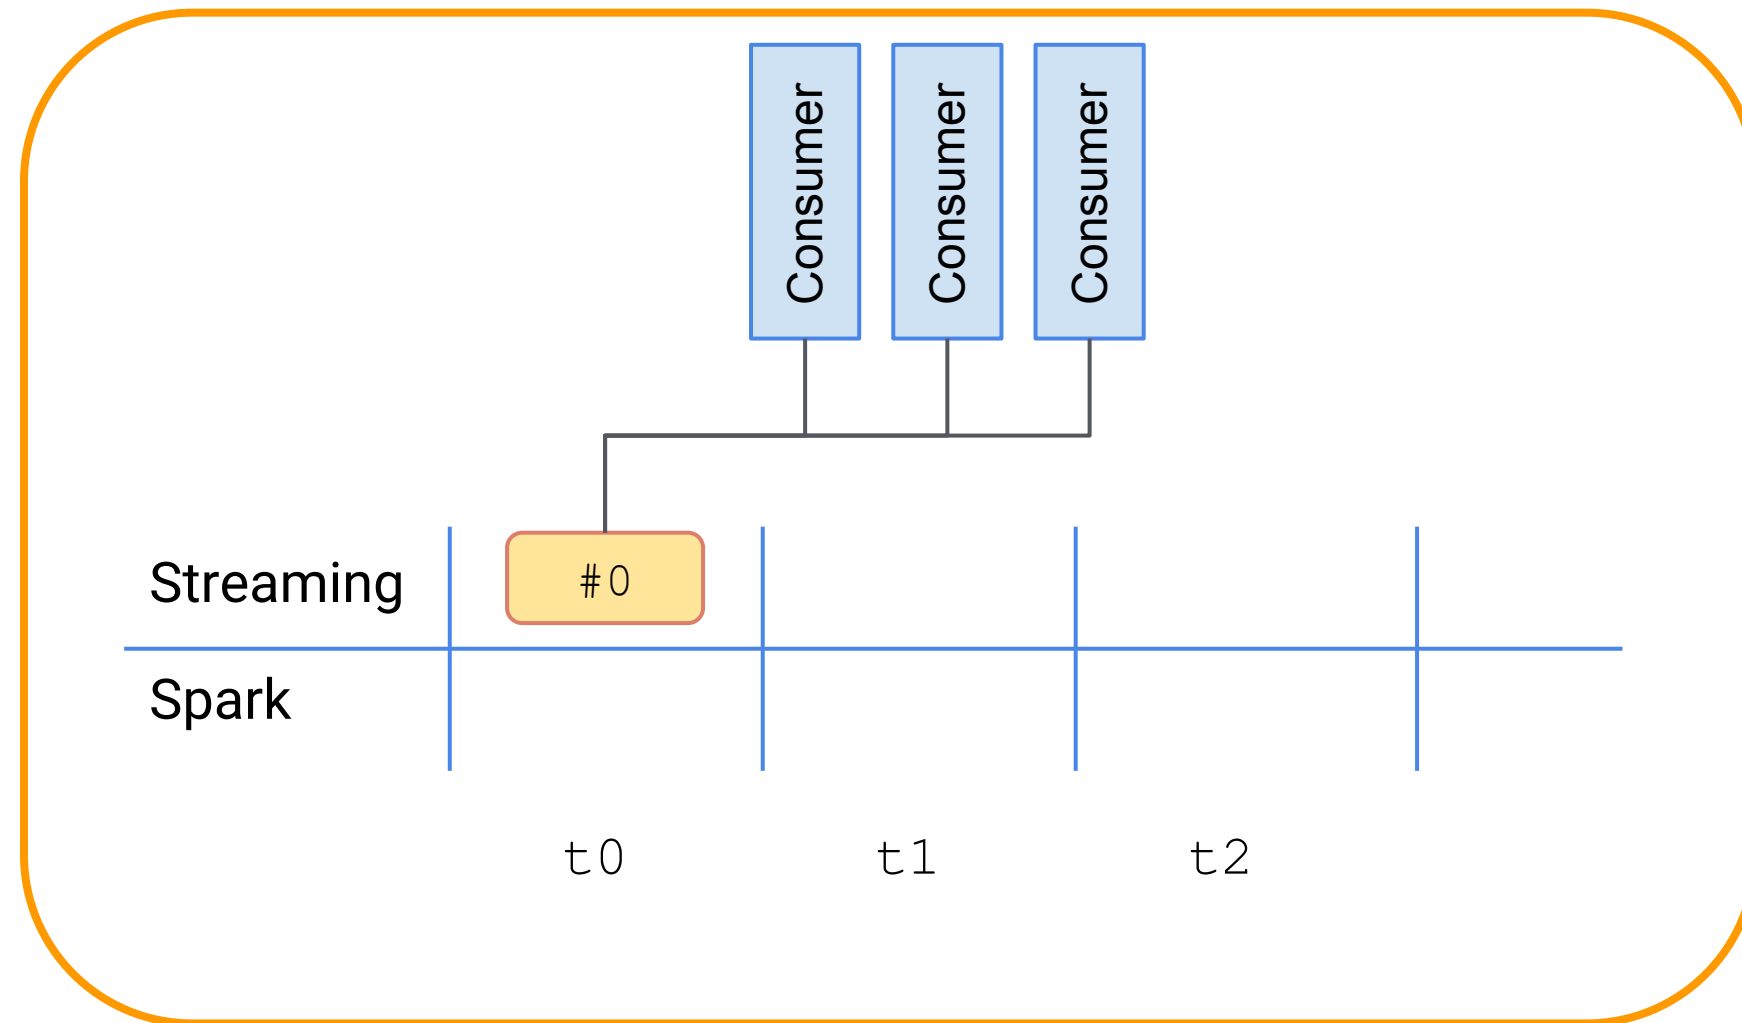
`spark.master=spark://host:port`

Using a Cluster Manager

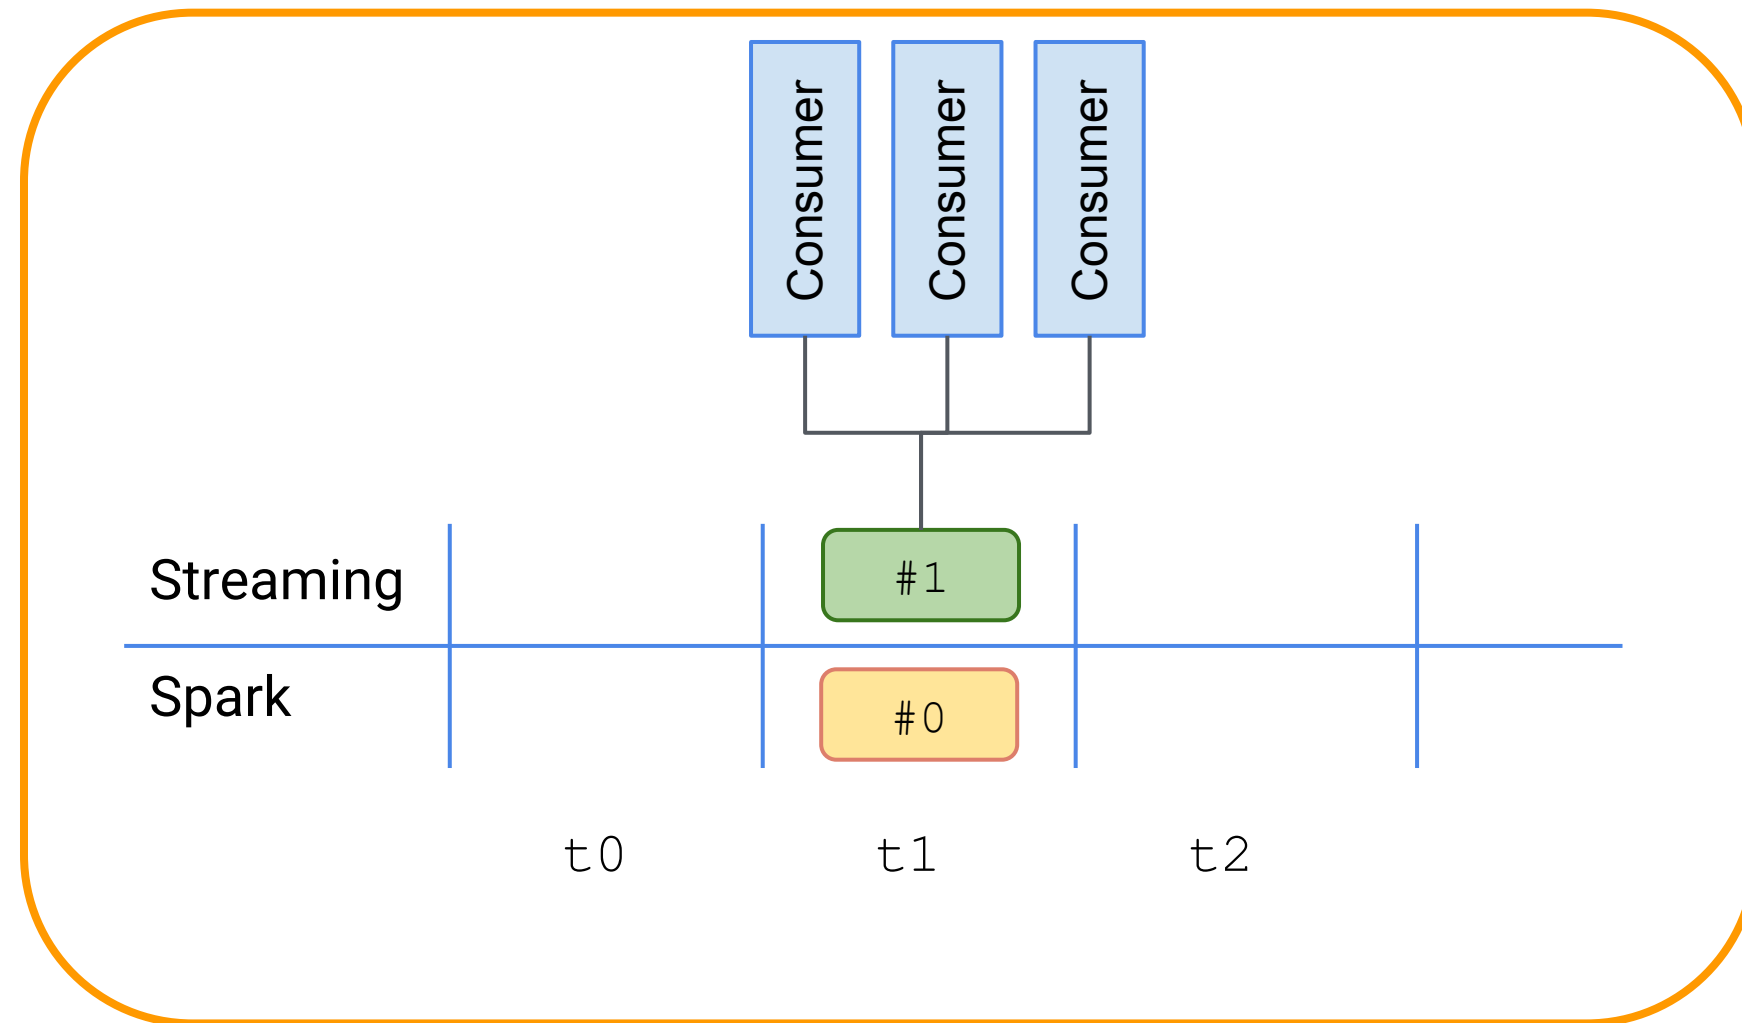


`spark.master=mesos://host:port`

# Scheduling

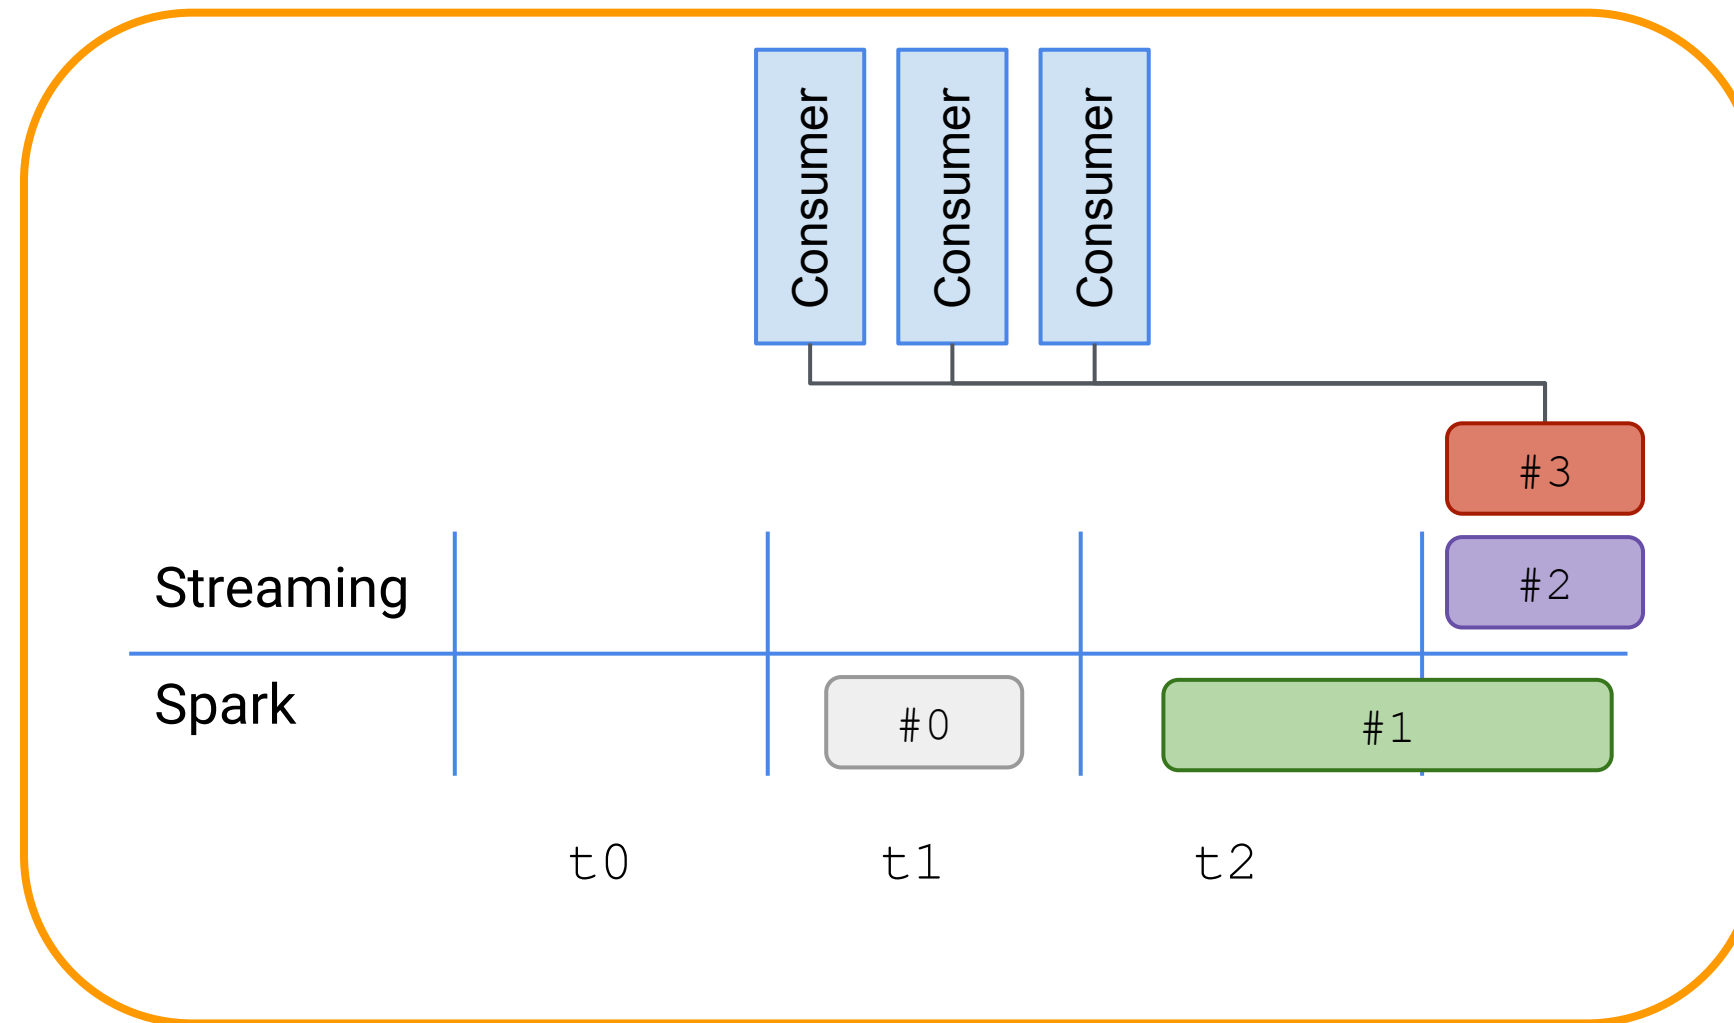


# Scheduling



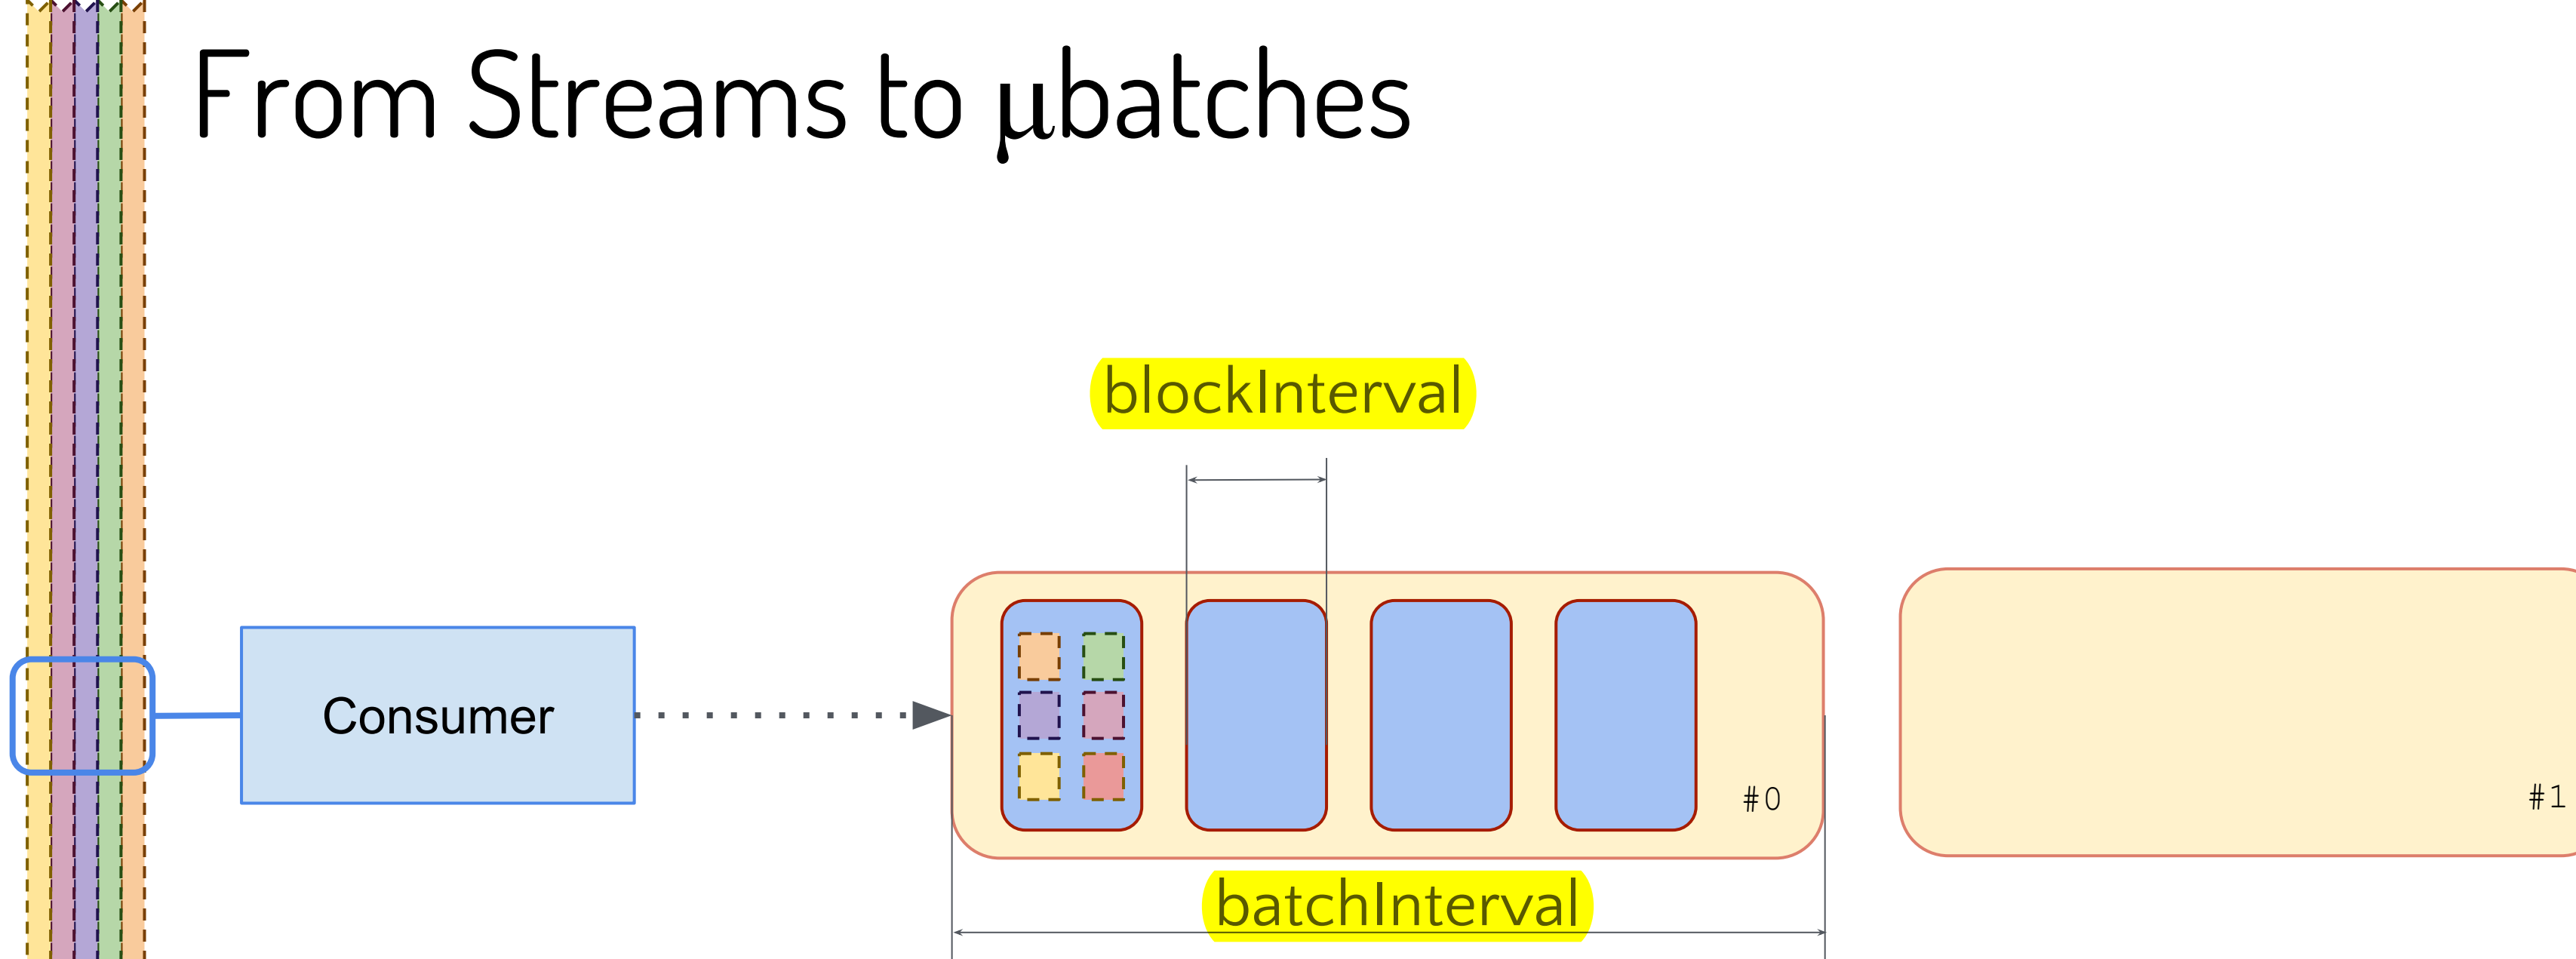
Process Time < Batch Interval

# Scheduling



Scheduling Delay


# From Streams to $\mu$ batches



Spark Streaming  
Spark

$$\text{\#partitions} = \text{receivers} \times \text{batchInterval} / \text{blockInterval}$$






# From Streams to $\mu$ batches


Spark Streaming

Spark


Partitions



RDD




Spark  
Executors




# From Streams to $\mu$ batches

Spark Streaming

Spark







# From Streams to $\mu$ batches

Spark Streaming

Spark







# From Streams to $\mu$ batches

Spark Streaming


Spark



# From Streams to $\mu$ batches

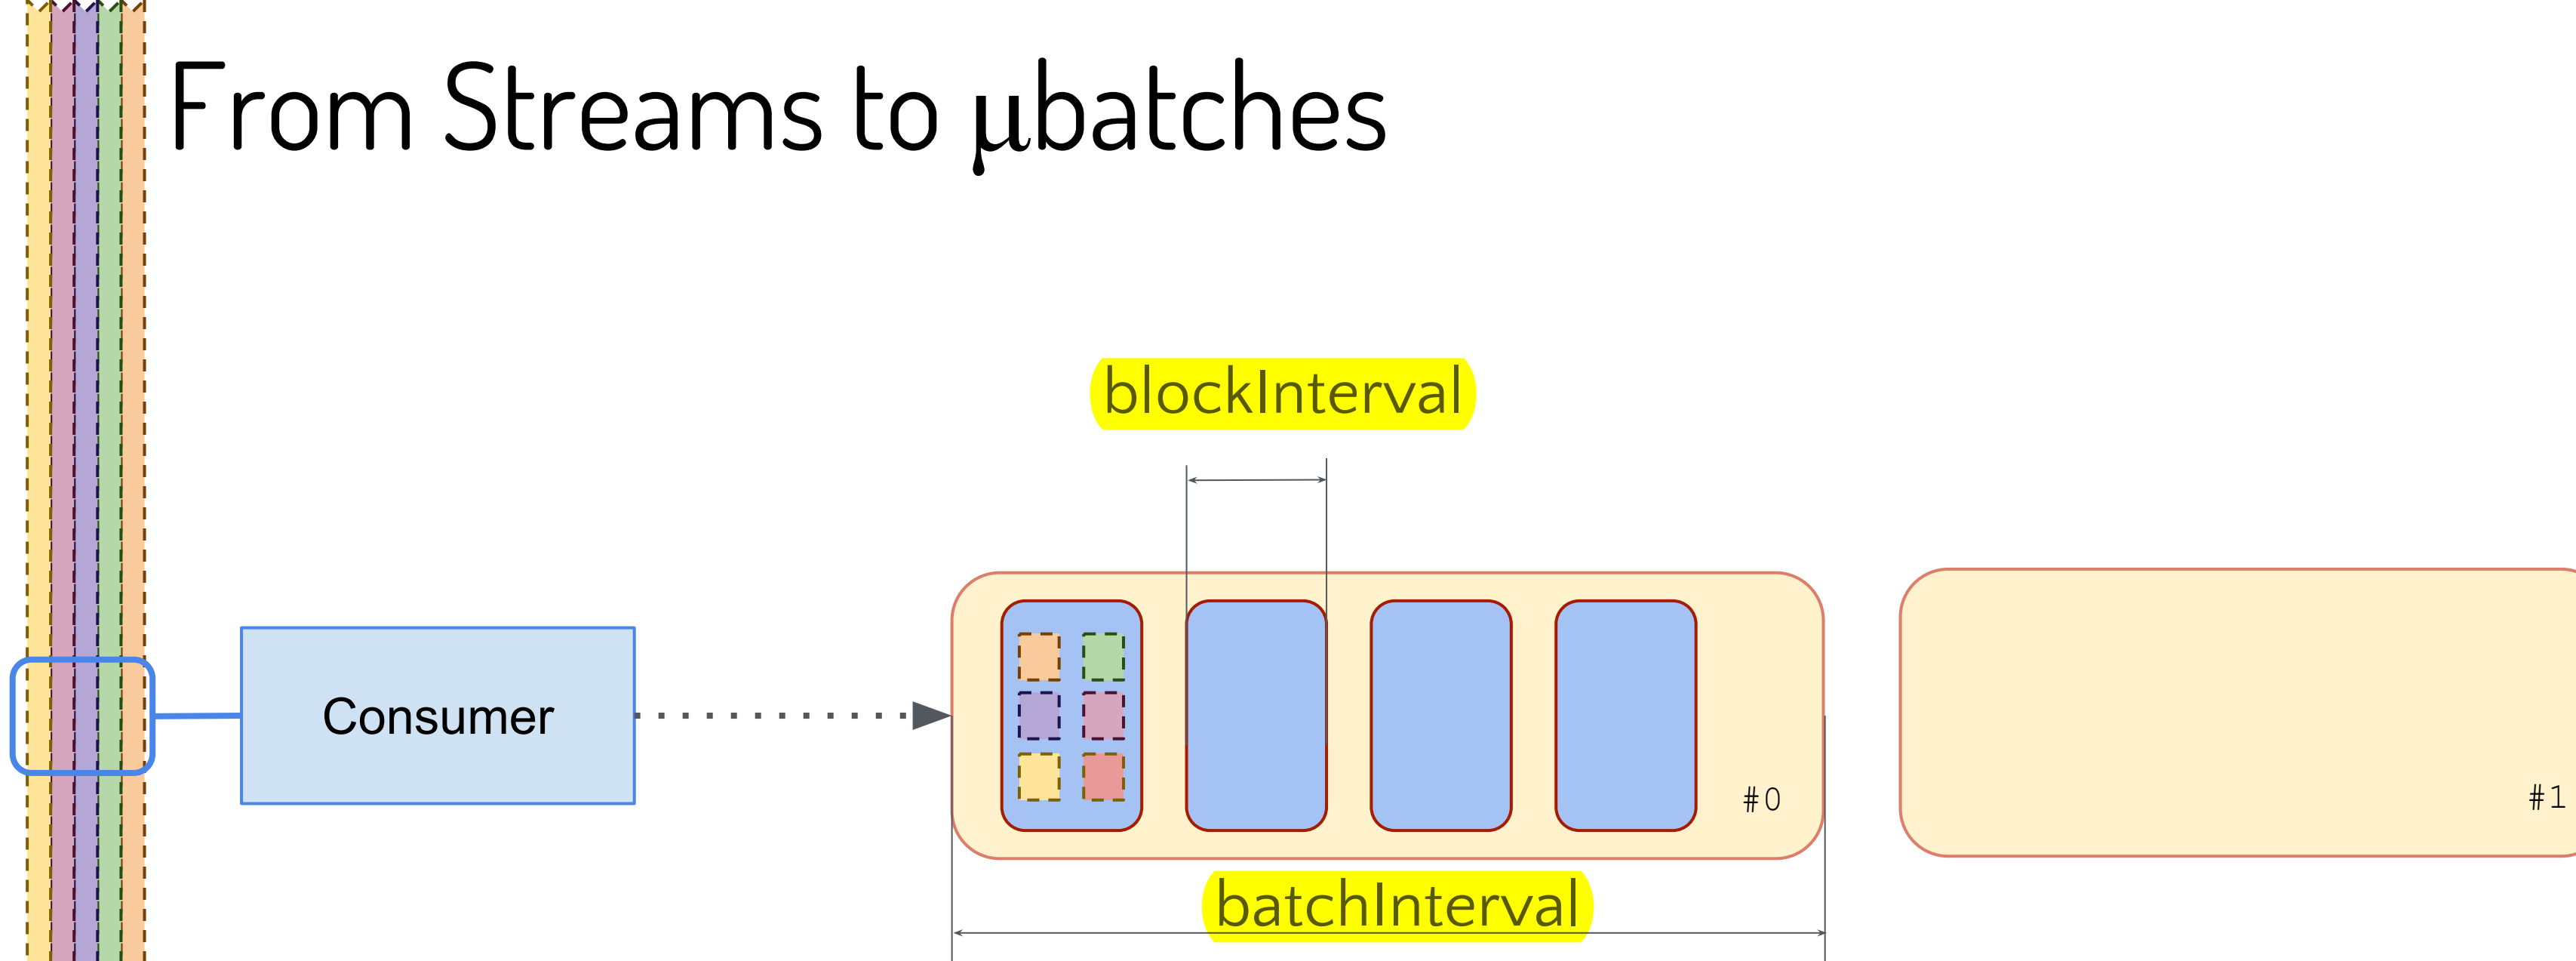


# From Streams to $\mu$ batches





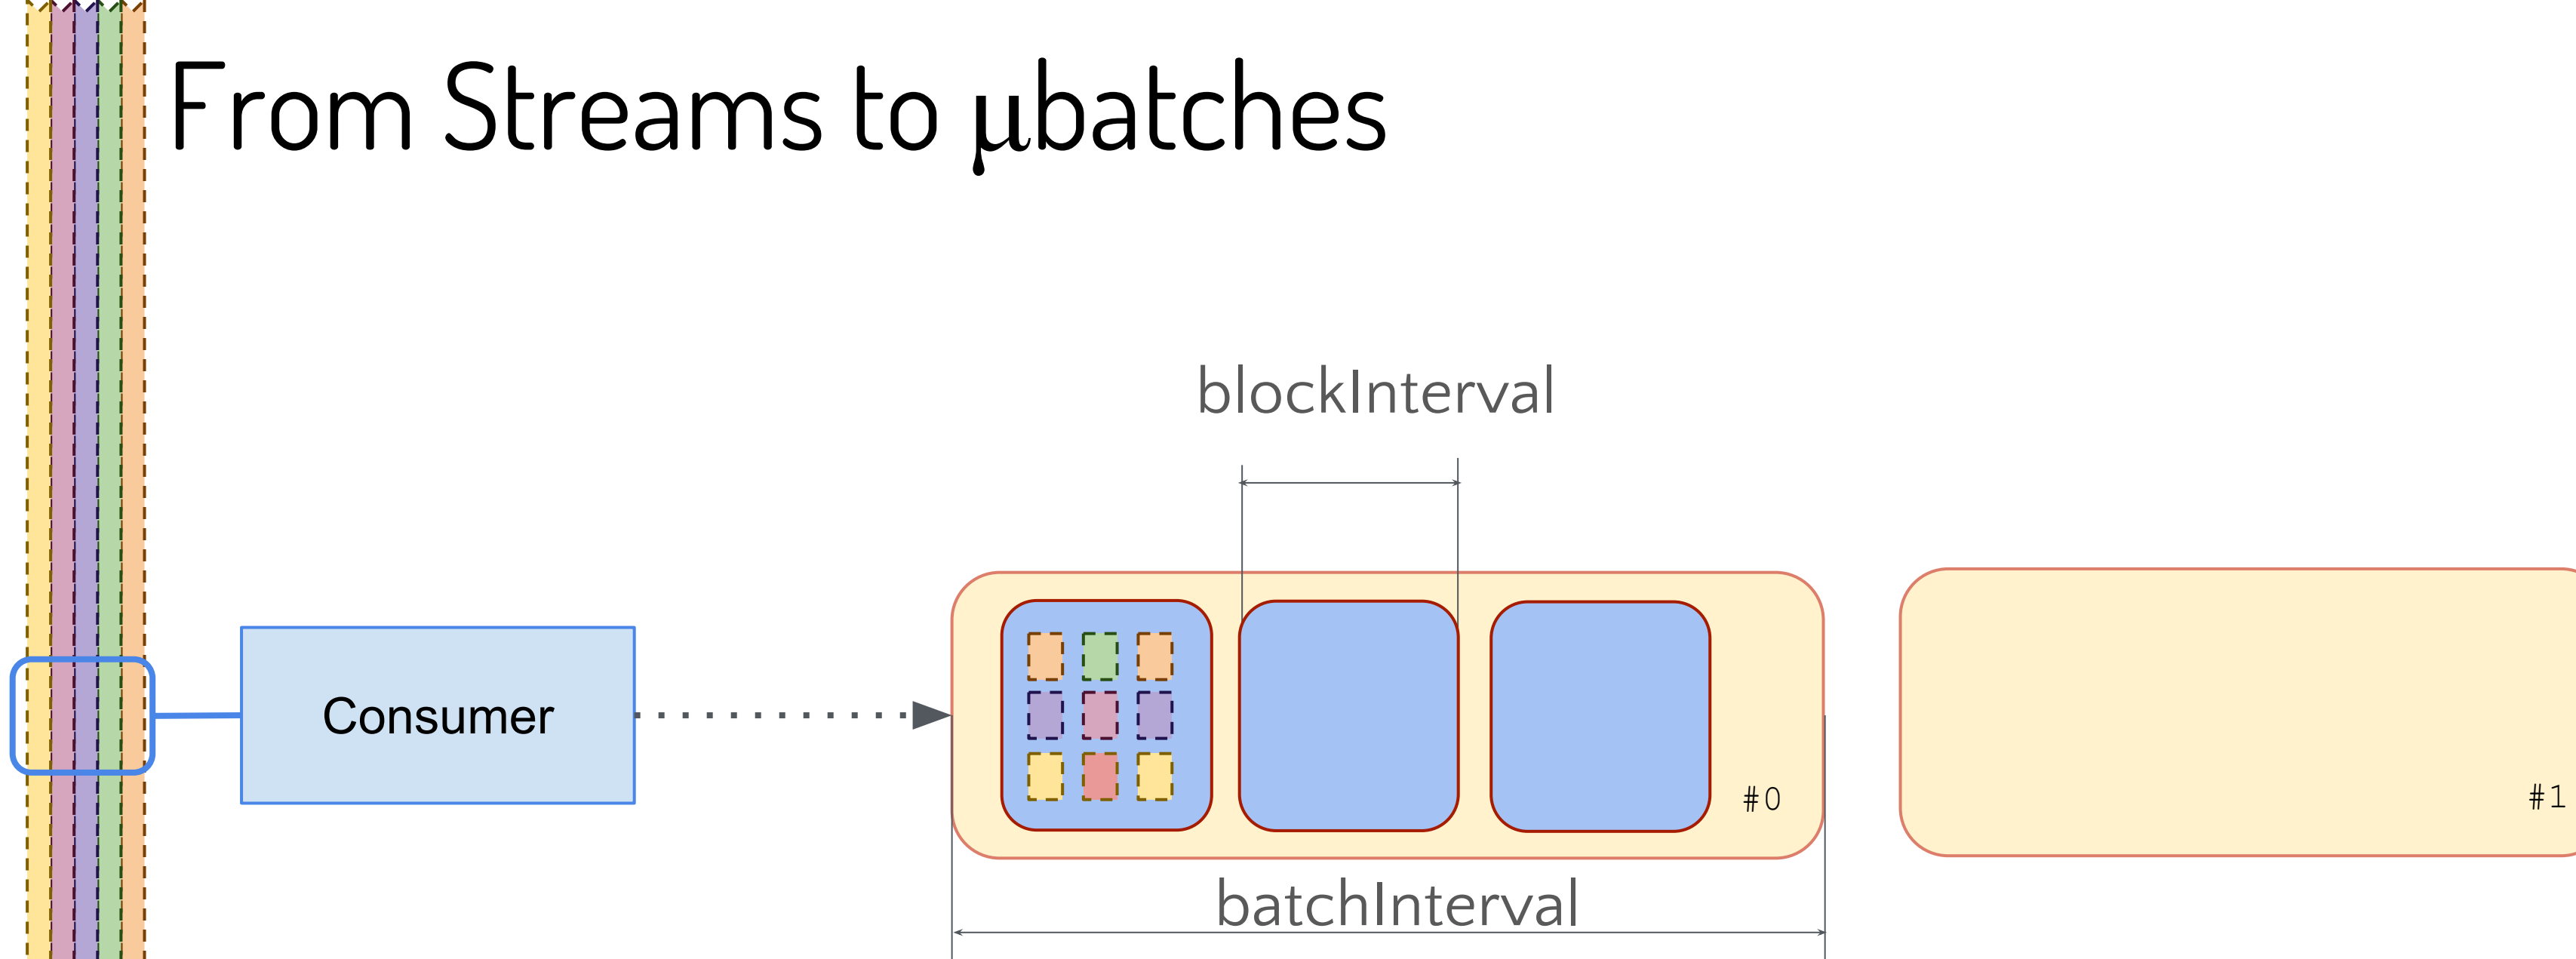
# From Streams to $\mu$ batches



Spark Streaming  
Spark

$$\text{\#partitions} = \text{receivers} \times \text{batchInterval} / \text{blockInterval}$$

# From Streams to $\mu$ batches




Spark Streaming

Spark

$$\text{spark.streaming.blockInterval} = \text{batchInterval} \times \frac{\text{receivers}}{(\text{partitionFactor} \times \text{sparkCores})}$$

# The Importance of Caching

```
dstream.foreachRDD { rdd =>  
  rdd.cache() // cache the RDD before iterating!  
  
  keys.foreach{ key =>  
    rdd.filter(elem=> key(elem) == key).saveAsFooBar(...)   
  }  
  
  rdd.unpersist()  
}
```




# The Receiver model

```
1  val ssc: StreamingContext = ???
2  val kafkaParams: Map[String, String] = Map("group.id" -> "terran", ...)
3  val readParallelism = 5
4  val topics = Map("zerg.hydra" -> 1)
5
6  val kafkaDStreams = (1 to readParallelism).map { _ =>
7      KafkaUtils.createStream(ssc, kafkaParams, topics, ...)
8  }
9  //> collection of five *input* DStreams = handled by five receivers/tasks
10
11  val unionDStream = ssc.union(kafkaDStreams) // often unnecessary, just show
12  //> single DStream
13
14  val processingParallelism = 20
15  val processingDStream = unionDStream(processingParallelism)
16  //> single DStream but now with 20 partitions
```

spark.streaming.receiver.maxRate  
Fault tolerance ? WAL

# Direct Kafka Stream



# Kafka: The Receiver-less model

- 😊 *Simplified Parallelism*
- 😊 *Efficiency*
- 😊 *Exactly-once semantics*
- Less degrees of freedom

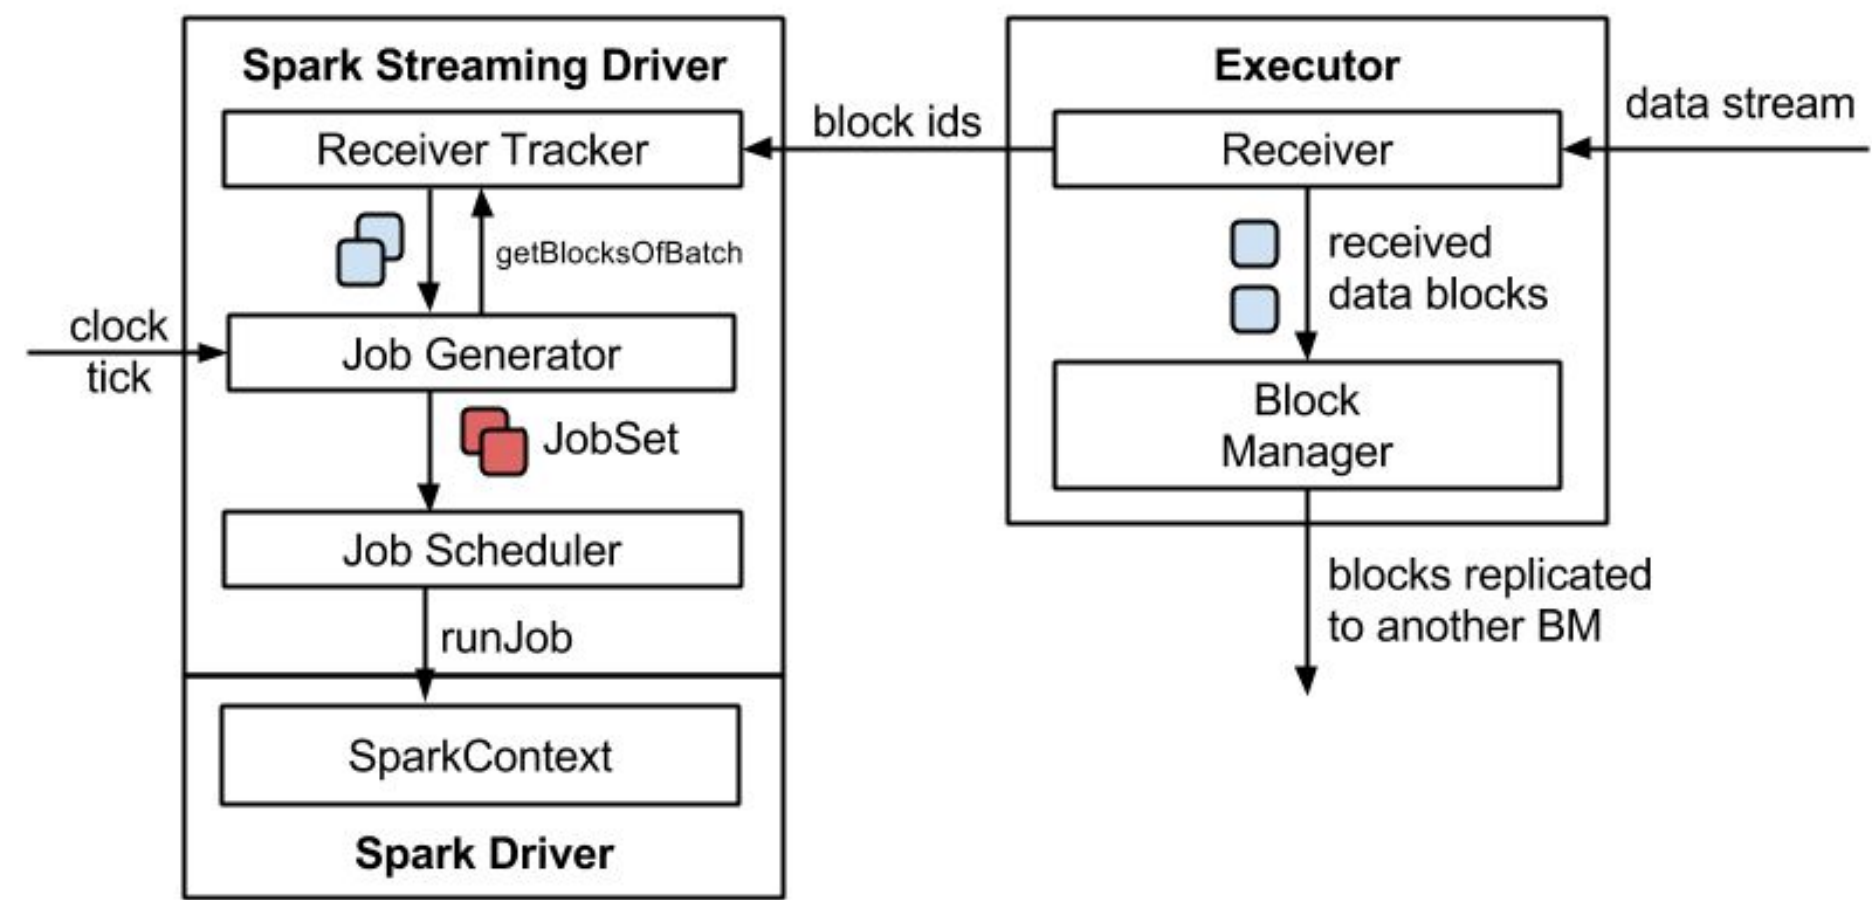
```
val directKafkaStream = KafkaUtils.  
createDirectStream[  
  [key class],  
  [value class],  
  [key decoder class],  
  [value decoder class] ](  
  streamingContext, [map of Kafka parameters], [set  
of topics to consume]  
)
```

`spark.streaming.kafka.maxRatePerPartition`



# Spark Streaming (v1.5) made Reactive

Backpressure support



proportional-integral-derivative controller (PID controller)





## Demo 2

# Spark Streaming Performance



# Resources

Spark Streaming Official Programming Guide:

<http://spark.apache.org/docs/latest/streaming-programming-guide.html>

Backpressure in Spark Streaming:

<http://blog.garillot.net/post/121183250481/a-quick-update-on-spark-streaming-work-since-i>

The Virdata's Spark Streaming tuning guide:

<http://www.virdata.com/tuning-spark/>

Spark Summit Presentations:

<https://spark-summit.org/>

Diving into Spark Streaming Execution Model:

<https://databricks.com/blog/2015/07/30/diving-into-spark-streamings-execution-model.html>

Kafka direct approach:

<https://github.com/koeninger/kafka-exactly-once/blob/master/blogpost.md>



Questions?



# Thanks!

**Gerard Maas**

  @maasg

[www.virdata.com](http://www.virdata.com)

- we're hiring -